

Vivek Sharma

8828482583 | vivektusharma@gmail.com | linkedin.com/in/vivek-sharma-64951b24b | github.com/Vivek02Sharma

EDUCATION

MCC, University of Mumbai

BSc in Information Technology; CGPA: 7.37

Mumbai, India

2022 – 2025

GVM College

High School (HSC Science); 65%

Mumbai, India

2020 – 2022

PROJECTS

MedVec-Scratch | Python, PyTorch, BPE Tokenizer, HuggingFace, Transformers

[Project Link]

- Built custom medical sentence embedding model from scratch using Siamese Transformer (4 layers, 256 dim) with custom BPE tokenizer (30k vocab), trained on 233k medical triplets using Triplet Margin Loss for healthcare semantic search.

Federal Registry RAG Agentic System | Python, Groq LLM, MySQL, Streamlit, Asyncio

[Project Link]

- Developed full-stack RAG system for US Federal Registry queries with Groq LLM, async pipelines, MySQL backend, and agentic tool calling architecture for real-time document retrieval.

Student Performance Analysis and Prediction | Python, Scikit-learn, XGBoost, Plotly, Streamlit

[Project Link]

- Built Streamlit dashboard for academic performance prediction using XGBoost with role-based access, trend analysis, and MongoDB backend for SGPA forecasting.

EXPERIENCE

Anvex AI Technologies Private Limited

September 2025 – Present

Jr. Software Engineer Machine Learning

Mumbai, India

- Built RAG system for AnvexSpeak, a telephony AI platform using Twilio/Plivo with multi-source data extraction (client databases, Google Drive, web scraping), BAAI/bge-large-en-v1.5 embeddings, and advanced filtering with hash256 for client data separation.
- Created custom LLM wrapper supporting open-source (Gemma3, Qwen, Mistral, Llama) and closed-source models (OpenAI) using LangChain/LangGraph; deployed with Qdrant (production) and ChromaDB (development) for real-time document-grounded responses.

Anvex AI Technologies Private Limited

June 2025 – August 2025

Machine Learning Intern

Mumbai, India

- Developed AVA, a RAG-powered policy chatbot using LangChain/LangGraph for query routing, Qdrant vector store, and SambaNova Llama-4 across 26+ policies with FastAPI endpoints achieving 800-1500ms response time and confidence scoring.

Techathon

February 2025

Mulund College of Commerce

Mumbai, India

- Built and trained machine learning models to predict equipment failures using sensor data, historical maintenance records, feature engineering, hyperparameter tuning, and integrated into backend using Flask/FastAPI for real-time predictive maintenance.

OPEN SOURCE CONTRIBUTION

Universal-Box

2024

Contributor (Data Science Template)

- Developed and added Data Science template with starter configurations for Jupyter, pandas, scikit-learn, structured workflows, and example notebooks, collaborating with maintainers to align with best practices.

TECHNICAL SKILLS

Languages: Python, SQL

Frameworks: TensorFlow, Flask, FastAPI, LangChain, LangGraph

Developer Tools: Docker, Linux, Git, GitHub, VS Code, Jupyter Lab, PyCharm

Libraries: PyTorch, Scikit-learn, pandas, NumPy, Matplotlib, Seaborn, Plotly

Vector Databases: Qdrant, ChromaDB